



# Generalized linear models, binary data

## Regression models

Brian Caffo, Jeff Leek and Roger Peng  
Johns Hopkins Bloomberg School of Public Health

# Key ideas

- Frequently we care about outcomes that have two values
  - Alive/dead
  - Win/loss
  - Success/Failure
  - etc
- Called binary, Bernoulli or 0/1 outcomes
- Collection of exchangeable binary outcomes for the same covariate data are called binomial outcomes.

# Example Baltimore Ravens win/loss

## Ravens Data

```
download.file("https://dl.dropboxusercontent.com/u/7710864/data/ravensData.rda"
, destfile="./data/ravensData.rda", method="curl")
load("./data/ravensData.rda")
head(ravensData)
```

	ravenWinNum	ravenWin	ravenScore	opponentScore
1	1	W	24	9
2	1	W	38	35
3	1	W	28	13
4	1	W	34	31
5	1	W	44	13
6	0	L	23	24

# Linear regression

$$RW_i = b_0 + b_1 RS_i + e_i$$

$RW_i$  - 1 if a Ravens win, 0 if not

$RS_i$  - Number of points Ravens scored

$b_0$  - probability of a Ravens win if they score 0 points

$b_1$  - increase in probability of a Ravens win for each additional point

$e_i$  - residual variation due

# Linear regression in R

```
lmRavens <- lm(ravensData$ravenWinNum ~ ravensData$ravenScore)
summary(lmRavens)$coef
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.2850	0.256643	1.111	0.28135
ravensData\$ravenScore	0.0159	0.009059	1.755	0.09625

# Odds

**Binary Outcome 0/1**

$$RW_i$$

**Probability (0,1)**

$$\Pr(RW_i | RS_i, b_0, b_1)$$

**Odds**  $(0, \infty)$

$$\frac{\Pr(RW_i | RS_i, b_0, b_1)}{1 - \Pr(RW_i | RS_i, b_0, b_1)}$$

**Log odds**  $(-\infty, \infty)$

$$\log \left( \frac{\Pr(RW_i | RS_i, b_0, b_1)}{1 - \Pr(RW_i | RS_i, b_0, b_1)} \right)$$

# Linear vs. logistic regression

## Linear

$$RW_i = b_0 + b_1 RS_i + e_i$$

or

$$E[RW_i | RS_i, b_0, b_1] = b_0 + b_1 RS_i$$

## Logistic

$$\Pr(RW_i | RS_i, b_0, b_1) = \frac{\exp(b_0 + b_1 RS_i)}{1 + \exp(b_0 + b_1 RS_i)}$$

or

$$\log\left(\frac{\Pr(RW_i | RS_i, b_0, b_1)}{1 - \Pr(RW_i | RS_i, b_0, b_1)}\right) = b_0 + b_1 RS_i$$

# Interpreting Logistic Regression

$$\log\left(\frac{\Pr(\text{RW}_i | \text{RS}_i, b_0, b_1)}{1 - \Pr(\text{RW}_i | \text{RS}_i, b_0, b_1)}\right) = b_0 + b_1 \text{RS}_i$$

$b_0$  - Log odds of a Ravens win if they score zero points

$b_1$  - Log odds ratio of win probability for each point scored (compared to zero points)

$\exp(b_1)$  - Odds ratio of win probability for each point scored (compared to zero points)

# Odds

- Imagine that you are playing a game where you flip a coin with success probability  $p$ .
- If it comes up heads, you win  $X$ . If it comes up tails, you lose  $Y$ .
- What should we set  $X$  and  $Y$  for the game to be fair?

$$E[earnings] = Xp - Y(1 - p) = 0$$

- Implies

$$\frac{Y}{X} = \frac{p}{1 - p}$$

- The odds can be said as "How much should you be willing to pay for a  $p$  probability of winning a dollar?"
  - (If  $p > 0.5$  you have to pay more if you lose than you get if you win.)
  - (If  $p < 0.5$  you have to pay less if you lose than you get if you win.)

# Visualizing fitting logistic regression curves

```
x <- seq(-10, 10, length = 1000)
manipulate(
  plot(x, exp(beta0 + betal * x) / (1 + exp(beta0 + betal * x)),
       type = "l", lwd = 3, frame = FALSE),
  betal = slider(-2, 2, step = .1, initial = 2),
  beta0 = slider(-2, 2, step = .1, initial = 0)
)
```

# Ravens logistic regression

```
logRegRavens <- glm(ravensData$ravenWinNum ~ ravensData$ravenScore,family="binomial")  
summary(logRegRavens)
```

Call:

```
glm(formula = ravensData$ravenWinNum ~ ravensData$ravenScore,  
family = "binomial")
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.758	-1.100	0.530	0.806	1.495

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-1.6800	1.5541	-1.08	0.28
ravensData\$ravenScore	0.1066	0.0667	1.60	0.11

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 24.435 on 19 degrees of freedom


Residual deviance: 20.895 on 18 degrees of freedom

AIC: 24.89

Number of Fisher Scoring iterations: 5

# Ravens fitted values

```
plot(ravensData$ravenScore,logRegRavens$fitted,pch=19,col="blue",xlab="Score",ylab="Prob Ravens Win")
```



# Odds ratios and confidence intervals

```
exp(logRegRavens$coeff)
```

```
(Intercept) ravensData$ravenScore  
0.1864      1.1125
```

```
exp(confint(logRegRavens))
```

```
2.5 % 97.5 %  
(Intercept) 0.005675 3.106  
ravensData$ravenScore 0.996230 1.303
```

# ANOVA for logistic regression

```
anova(logRegRavens,test="Chisq")
```

Analysis of Deviance Table

Model: binomial, link: logit

Response: ravensData\$ravenWinNum

Terms added sequentially (first to last)

	Df	Deviance	Resid. Df	Resid. Dev	Pr(>Chi)
NULL			19	24.4	
ravensData\$ravenScore	1	3.54	18	20.9	0.06 .
<hr/>					
Signif. codes:	0	****	0.001	***	0.01 **
					0.05 .
					0.1 ' '
					1

# Interpreting Odds Ratios

- Not probabilities
- Odds ratio of 1 = no difference in odds
- Log odds ratio of 0 = no difference in odds
- Odds ratio < 0.5 or > 2 commonly a "moderate effect"
- Relative risk  $\frac{\Pr(\text{RW}_i | \text{RS}_i=1)}{\Pr(\text{RW}_i | \text{RS}_i=0)}$  often easier to interpret, harder to estimate
- For small probabilities RR  $\approx$  OR but **they are not the same!**

[Wikipedia on Odds Ratio](#)

# Further resources

- [Wikipedia on Logistic Regression](#)
- [Logistic regression and glms in R](#)
- Brian Caffo's lecture notes on: [Simpson's paradox](#), [Case-control studies](#)
- [Open Intro Chapter on Logistic Regression](#)